



Misspecifying the Shape of a Random Effects Distribution: Why Getting It Wrong May Not Matter

Author(s): Charles E. McCulloch and John M. Neuhaus

Source: *Statistical Science*, Vol. 26, No. 3 (August 2011), pp. 388-402

Published by: Institute of Mathematical Statistics

Stable URL: <http://www.jstor.org/stable/23059138>

Accessed: 09-10-2017 16:51 UTC

REFERENCES

Linked references are available on JSTOR for this article:

http://www.jstor.org/stable/23059138?seq=1&cid=pdf-reference#references_tab_contents

You may need to log in to JSTOR to access the linked references.

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <http://about.jstor.org/terms>



Institute of Mathematical Statistics is collaborating with JSTOR to digitize, preserve and extend access to *Statistical Science*

Misspecifying the Shape of a Random Effects Distribution: Why Getting It Wrong May Not Matter

Charles E. McCulloch and John M. Neuhaus

Abstract. Statistical models that include random effects are commonly used to analyze longitudinal and correlated data, often with strong and parametric assumptions about the random effects distribution. There is marked disagreement in the literature as to whether such parametric assumptions are important or innocuous. In the context of generalized linear mixed models used to analyze clustered or longitudinal data, we examine the impact of random effects distribution misspecification on a variety of inferences, including prediction, inference about covariate effects, prediction of random effects and estimation of random effects variances. We describe examples, theoretical calculations and simulations to elucidate situations in which the specification is and is not important. A key conclusion is the large degree of robustness of maximum likelihood for a wide variety of commonly encountered situations.

Key words and phrases: Maximum likelihood, mixed models, parametric modeling.

1. INTRODUCTION

Statistical models that include random effects are commonly used to analyze longitudinal and clustered data. When generalized linear mixed models (McCulloch, Searle and Neuhaus, 2008) and maximum likelihood estimation are used to analyze such data, strong, parametric assumptions about the random effects distribution are typically made. Are inferences sensitive to this specification?

One body of research has indicated that the impact of misspecification of the *shape* of the random effects distribution is slight, especially for estimating regression parameters other than the intercept (e.g., Neuhaus, Hauck and Kalbfleisch, 1992; Neuhaus, Kalbfleisch and Hauck, 1994; Butler and Louis, 1992; Heagerty and Kurland, 2001).

However, a number of authors have claimed sensitivity to parametric specification of a random effects distribution. An oft-quoted article is that of Heckman and Singer (1984), which (in describing performance of maximum likelihood estimation) states, “Both theoretical and empirical examples indicate that estimates of structural parameters obtained from conventional procedures are very sensitive to the choice of mixing distributions.” Other authors have suggested more flexible distributional assumptions for the random effects to protect against misspecification. Approaches include nonparametric maximum likelihood (Aitkin, 1999; Agresti, Caffo and Ohman-Strickland, 2004), more flexible parametric distributions (Zhang et al., 2008), marginalized mixed effects models (Heagerty and Zeger, 2000), h-likelihood approaches that can be easily adapted to fit different distributions (Lee and Nelder, 2004), families of parametric distributions (Piepho and McCulloch, 2004), mixtures of normal distributions (Lesaffre and Molenberghs, 2001) and “smooth” nonparametric fits (Zhang and Davidian, 2001). A recent article (Huang, 2009) encapsulates the sentiment used to justify these approaches:

Charles E. McCulloch is Professor and Head, Department of Epidemiology and Biostatistics, 185 Berry Street, San Francisco, California 94107, USA (e-mail: chuck@biostat.ucsf.edu). John M. Neuhaus is Professor, Department of Epidemiology and Biostatistics, 185 Berry Street, San Francisco, California 94107, USA (e-mail: john@biostat.ucsf.edu).

For computational convenience, random effects in GLMMs are almost routinely assumed to be normal. However, the normality assumption may be unrealistic in some applications. ... Early investigation to address this concern suggested that misspecifying the models for the random effects usually only results in a small amount of bias in the maximum likelihood estimators (MLEs) for the fixed effects (Neuhaus, Hauck and Kalbfleisch, 1992). However, more recently, many authors have found that likelihood-based inference can be severely affected if the random-effect model is misspecified. For example, Heagerty and Kurland (2001) computed the asymptotic bias in the MLEs for the parameters in a logistic mixed model in four instances of random-effect model misspecification. They concluded that incorrect assumptions on the random effects can lead to substantial bias in the MLEs for the fixed effects. Agresti, Caffo and Ohman-Strickland (2004) conducted empirical studies on the impact of model misspecification for the random effects in GLMMs, showing that the MLEs for the fixed effects can be very sensitive to the assumed random effect model. Finally, Litière, Alonso and Molenberghs (2007) used simulation to show that the type I and type II errors of tests for the mean structure in a logistic mixed model can be seriously affected by violations of the random-effect model.

A body of work with a slightly different focus has been that of estimating the shape of the random effects distribution, also by hypothesizing more flexible distributional fits for the random effects. Though this is different from assessing misspecification, it is a closely related problem and those methods may help diagnose misspecification in the situations in which sensitivity to misspecification is of concern. Approaches to this problem have used mixtures of normal distributions (Magder and Zeger, 1996; Caffo, An and Rohde, 2007), and nonparametric and “smooth” nonparametric fits (Laird, 1978; Davidian and Gallant, 1993; Zhang and Davidian, 2001; Ghidey, Lesaffre and Eilers, 2004). See Ghidey, Lesaffre and Verbeke (2010) for a recent review article in the context of linear mixed models.

We use data from the Heart and Estrogen Replacement Study (HERS) (Hulley et al., 1998) to illustrate some of our findings. HERS was a randomized, blinded, placebo controlled trial for women with previous coronary disease. The study enrolled 2,763 women and followed them annually for five subsequent visits, generating longitudinal data. More detail is given in Section 11, but we foreshadow the results here. We fit logistic regression models with random intercepts to model whether or not a woman had high blood pressure as a function of her body mass index, whether she was on hypertensive medication, and a trend over the visits. For the random intercepts the models assumed one of four distributions: normal, a centered and scaled exponential distribution, a Tukey(g, h) distribution or a 3 point discrete distribution, quite different parametric assumptions. The regression parameters were very similar, despite evidence for differences in overall model fit. How can we reconcile this with the claims above?

We will argue that concerns over the parametric specification of random effects distributions are sometimes valid, but oftentimes misplaced. This is due to three reasons: (1) sensitivities are restricted to aspects of the estimation that are not typically of interest, (2) the situations considered are unfairly extreme, or (3) that close scrutiny of published results does not actually support sensitivity to misspecification.

Our paper is organized as follows. In the next section we outline a series of examples and their associated inferences. We then describe the basic statistical model. The sections that follow look at the impact of misspecification on various aspects of the model. The remaining sections consider some simulation studies, the HERS example (in more detail), some remarks on more complicated random effects structures, and we conclude with a brief summary.

2. INFERENCE SETTINGS

A key argument we make is that the robustness is dependent on the inferential setting. Perhaps the most common inference for which a regression model is used is estimation and testing of regression coefficients associated with a covariate. In the clustered data setting, a useful distinction is whether a covariate is a “between-cluster” covariate, meaning that it is constant over the units in a cluster or a “within-cluster” covariate, meaning that it varies within a cluster, but has an average that is constant between clusters. For example, in a longitudinal study with participants being clusters and all participants contributing data across

4 time points, the ethnicity of the subject would be a between-cluster covariate and the visit number would be a within-cluster covariate. Of course, covariates are often neither purely between or within, for example, the body mass index of a participant measured over time, in which case it is often useful for pedagogical, substantive or technical reasons to decompose a covariate into its purely-between and purely-within components (Neuhaus and Kalbfleisch, 1998; Raudenbush and Bryk, 2002; Neuhaus and McCulloch, 2006). Other inferential goals include using the fitted model to generate predictions, generating predicted values of the random components of the model, or estimating aspects of the distribution of the random components of the model, for example, as in a variance components analysis (Searle, Casella and McCulloch, 1992).

2.1 Examples

To illustrate these distinctions, we begin with a series of examples.

2.1.1 Estimate the effect of a within-cluster covariate. Metlay et al. (2007) evaluated the effectiveness of an educational program that attempted to reduce the inappropriate use of antibiotics for conditions nonresponsive to antibiotics. The study was a cluster randomized trial, randomized at the level of the hospital, with 8 intervention and 8 control hospitals. Measurements were taken at each hospital in the year before and after the interventions were introduced. The data were further clustered by physician within hospital. The primary inference involved a comparison of responses before and after the intervention, a within-cluster (= hospital) covariate setting.

2.1.2 Estimate the effect of a between-cluster covariate. In the previous example, 8 of the hospitals were Veteran's Administration (VA) hospitals and 8 were non-VA. A secondary goal was to compare the rates of inappropriate usage between VA and non-VA hospitals, a between-cluster covariate setting.

2.1.3 Derive predictions based on the fixed effects. Auble et al. (2007) studied the performance of four clinical prediction rules for estimating the risk of death or serious complications for patients admitted to the hospital with a diagnosis of heart failure. The goal was to develop a rule that could stratify patients into low and high risk groups. One of the models used was a random effects logistic regression model accommodating clustering of outcomes by hospital. The inferential goal was to develop a prediction rule based on the fixed effects in the model.

2.1.4 Predict random effects. Zhang et al. (2008) wanted to estimate subject-specific rates of disease progression in chronic kidney disease patients. They estimated subject-specific slopes of the change in glomerular filtration rate (a measure of kidney function) over time. They used a linear mixed model with clustering by subject and non-normally distributed random effects. The inferential goal was to predict the realized values of the random effects.

2.1.5 Estimate variance components. Selby et al. (1996) studied the variation in the rates of coronary angiography (a diagnostic procedure using X-rays and a special dye to diagnose coronary artery disease) across 16 hospitals. Significant variation after adjusting for differences in patient populations indicates over- or under-use of the procedure, leading to poor patient outcomes and/or wasted resources. The inferential goal was to estimate the variance of the hospital random effects after adjusting for patient characteristics.

2.2 Longitudinal and Clustered Studies

Each of the above examples describes clustered or longitudinal studies, in which repeated observations of the outcome are taken within clusters, for example, multiple measurements of kidney function on a patient. These are the types of scenarios we consider.

An immediate distinction to make is the different scenario considered by Heckman and Singer (1984), in which there is only a single observation of the outcome per cluster (in their case, a time-to-event outcome). In the normal linear mixed model setting, with only a single observation per cluster, the random effect and error term are completely confounded, leading to lack of identifiability. The presence of identifiability in the Heckman and Singer situation arises only through strong parametric assumptions, so it is not surprising that results are sensitive to the choice of distribution. Such settings do not represent longitudinal or clustered studies and we do not consider them further.

3. A GENERALIZED LINEAR MIXED MODEL

The model we will use to assess the impact of misspecification is a generalized linear mixed model for clustered data with random intercepts, b_i . There are fewer results for more complicated covariance structures, such as those generated by random intercept and slope models, which we discuss briefly in Section 12. Let Y_{it} represent the t th observation ($t = 1, \dots, n_i$)

within cluster i ($i = 1, \dots, m$). We assume that, conditional on the random effects, the Y_{it} are independent:

$$\begin{aligned} Y_{it}|b_i &\sim \text{independent } F_Y, \\ i &= 1, \dots, m; t = 1, \dots, n_i, \\ (1) \quad g(E[Y_{it}|b_i]) &= b_i + \mathbf{x}'_{it}\boldsymbol{\beta}, \\ b_i &\sim \text{i.i.d. } F_b, \\ E[b_i] &= 0 \quad \text{and} \quad \text{var}(b_i) = \sigma_b^2, \end{aligned}$$

where $g(\cdot)$ is a known link function, $\boldsymbol{\beta}$ is the parameter vector for the fixed effects, and \mathbf{x}_{it} is a vector of covariates for cluster i at time t . We consider the performance of maximum likelihood for estimating the parameters. We focus especially on the situation where the assumed distribution is normal, since this is the most commonly implemented choice in popular software packages such as SAS, Stata and R.

The model, (1), contains a number of specifications. The main specification we will consider is the *shape* of the random effects distribution, that is, the parametric form of the distribution, F_b , may be incorrectly specified.

We have briefer comments on two other aspects of the specification. We may be concerned that:

- Basic characteristics of the random effects distribution may depend on a covariate. For example, the mean or variance of F_b depends on a covariate.
- There is a dependence of the random effects distribution on cluster sample size. For example, the mean of F_b depends on n_i .

When the mean of the random effects distribution depends on a covariate, a fundamental relationship is introduced between the covariate and the distribution, potentially creating a serious bias in estimating the form of the relationship between the covariate and the outcome. Neuhaus and McCulloch (2006) discuss reasons for this, suggest decomposing covariates of interest into within- and between-cluster components and investigate conditions under which this decomposition can remove or reduce bias. Heagerty and Kurland (2001) study situations where the variance of the random effects distribution depends on a covariate and consider use of conditionally specified models (as we do here), as well as models that specify a marginal regression equation (e.g., logistic regression for binary outcomes) but incorporate random effects in an underlying conditional model. They show that, for estimating the conditionally specified parameters of (1), the

impact of highly unequal variances can lead to substantial bias. So both of these aspects of the specification of (1) are potentially important.

Some authors have argued (Hoffman, Sen and Weinberg, 2001; Williamson, Datta and Satten, 2003) that when the cluster sample size is incorrectly assumed to be independent of the random effects distribution, serious consequences may result. Neuhaus and McCulloch (2011) argue that this type of misspecification is really a variation of incorrect assumptions about the mixing distribution. If one assumes that the parameters in the joint distribution of n_i and X_{ij} are not functionally related to the parameters of interest, namely, those in the distribution of outcomes or random effects, then as a function of β , Neuhaus and McCulloch (2011) show that

$$\begin{aligned} f(\mathbf{y}_i, n_i, \mathbf{X}_i) \\ (2) \quad &\propto \int_b \left\{ \prod_{j=1}^{n_i} f_Y(y_{ij}|n_i, X_{ij}, b) \right\} dF_b(b|n_i, \mathbf{X}_i) \\ &= f(\mathbf{y}_i|n_i, \mathbf{X}_i). \end{aligned}$$

This is a useful representation of the joint likelihood of \mathbf{y}_i , n_i and \mathbf{X}_i since analysts would typically model the distribution of the responses conditional on covariates and sample size. In particular, ignoring the association of cluster size with response, one would base inference on the incorrect likelihood built up of terms

$$(3) \quad f^*(\mathbf{y}_i|\mathbf{X}_i, n_i) = \int_b \prod_{j=1}^{n_i} f_Y(y_{ij}|X_{ij}, b) dF_b^*(b),$$

where the asterisks denote assumed distributions.

Comparing equation (3), the assumed likelihood and (2), the true likelihood, we see that we can view (3) as arising from (2) but with a misspecified random effects distribution, namely, the conditional distribution of b given n_i and \mathbf{X}_i is incorrectly specified as $f_b^*(b)$. This is important because some of the comments we make below for random intercepts models concerning misspecification of the *shape* of the random effects distribution will also apply to incorrectly assuming that the cluster sample size is independent of the random intercept distribution. Situations in which the informative cluster size is related to the random slopes are more involved.

We next consider in more detail the impact of shape misspecification of the random effects distribution, organized by the inferential targets.

4. ESTIMATE/TEST A WITHIN-CLUSTER COVARIATE

Virtually every study of the impacts of misspecification has shown little impact on within-cluster covariates. Neuhaus, Hauck and Kalbfleisch (1992), Neuhaus, Kalbfleisch and Hauck (1994) and Heagerty and Kurland (2001) developed analytic results to assess the impact of misspecification of the random effects distribution using the theory of inference under misspecified models (Huber, 1967; Akaike, 1973; White, 1994). This theory shows that estimators obtained by maximizing the likelihood based on a misspecified random effects distribution converge to values that minimize the Kullback–Leibler divergence (Kullback, 1959) between the correctly specified and misspecified models. In cases such as the linear mixed effects model and binary matched pairs (Neuhaus, Kalbfleisch and Hauck, 1994), one can obtain closed form solutions for the values that minimize the Kullback–Leibler divergence. In these cases, the closed form solution shows that the estimator based on misspecified random intercepts is identical to the standard conditional likelihood estimator, an estimator unaffected by the random effects distribution. In other cases, one can show consistent estimation at $\beta = 0$ for generalized linear mixed models with misspecified random intercept distributions. For example, Neuhaus, Hauck and Kalbfleisch (1992) show consistent estimation at $\beta = 0$ for logistic models with misspecified random intercept distributions and Neuhaus and McCulloch (2011) extend the result to the entire class of generalized linear mixed models. In addition, one can obtain approximate solutions to Kullback–Leibler divergence minimizers using Taylor expansions (Neuhaus, Hauck and Kalbfleisch, 1992). Using this device, Neuhaus, Hauck and Kalbfleisch (1992) show little asymptotic bias in a logistic model with misspecified random intercepts. Heagerty and Kurland (2001), in their Table 1, show virtually no impact on the asymptotic bias in logistic regression models with a wide variety of gamma-distributed random effects. Chen, Zhang and Davidian (2002), in investigating bias and efficiency in logistic models, state that:

Estimation of β_2 , which corresponds to a covariate changing within individuals, suffers no loss of efficiency under misspecification of the random effects distribution. Similar results have been reported by Tao et al. (1999) and Zhang and Davidian (2001). ...a within individual covariate such as time is roughly ‘orthogonal’ to

among-individual effects. Thus, estimation of the associated regression coefficient may be less affected.

The above cited Zhang and Davidian article, investigating a linear mixed model, shows little impact on bias or efficiency.

So for the inferential goal which is arguably the most relevant to clustered or longitudinal data, misspecification of the shape of the random effects distribution has little or no effect.

5. ESTIMATE/TEST A BETWEEN-CLUSTER COVARIATE

Between-cluster covariates might be expected to be more sensitive to shape specification. We agree with Chen, Zhang and Davidian (2002), who state, “We conjecture that, because a cluster-level covariate such as treatment and the latent random effects both pertain to among-individual variation, misspecification of the random effects distribution would compromise quality of estimation of the corresponding regression coefficient.” However, the results of Neuhaus, Hauck and Kalbfleisch (1992) and Neuhaus and McCulloch (2011) on consistent estimation at $\beta = 0$ and the Taylor approximations of Neuhaus, Hauck and Kalbfleisch (1992) apply equally to within- and between-cluster covariates. In addition, the simulation results of Heagerty and Kurland (2001) show virtually no bias in estimates of between-cluster covariate effects.

Via simulation, Agresti, Caffo and Ohman-Strickland (2004), in their Table 3 and for a logistic model, showed moderate loss in efficiency when comparing assumed normal and two-point discrete distributions versus the same two true distributions with a small number of clusters (10). For a linear mixed model, Magder and Zeger (1996), in their Table 3, showed little or no loss in efficiency when assuming a normal random effects distribution when the distribution was actually skewed or bimodal, but moderate efficiency loss when the true distribution was a two-point discrete distribution.

Of more concern are the results of Zhang and Davidian (2001), Litière, Alonso and Molenberghs (2007) and Litière, Alonso and Molenberghs (2008), who investigated performance under highly non-normal distributions, but not as extreme as a two-point discrete distribution. Zhang and Davidian (2001) present simulation results for a linear mixed model with a true mixture of normals distribution but fit with an assumed normal distribution and show loss of

efficiency. Litière, Alonso and Molenberghs (2007) present simulation results for a logistic mixed model under a true “Power function” distribution (a continuous, but bounded range distribution) and assumed normal distribution and claimed a significant loss of power under an assumed normal distribution.

The scenario simulated by Zhang and Davidian was a balanced, linear mixed model, comparing clustered responses between two levels of a binary between-cluster covariate. The usual maximum likelihood based analysis estimates the between group difference with the ordinary least squares estimate (Searle, Casella and McCulloch, 1992), the difference in the group means. So the comparison simulated was essentially the performance of the t-test under non-normality. While the two-group t-test is well known to be highly robust (e.g., Rasch and Guiard, 2004), it is not impervious to non-normality and demonstrates loss of efficiency in this extreme violation of the normality assumption. However, the estimates are exactly unbiased.

The Litière, Alonso and Molenberghs (2007) and Litière, Alonso and Molenberghs (2008) simulations investigate the performance of tests that assume a normal distribution for the random effects while varying the true distributions. This is useful for understanding the robustness of methods available in standard software (which often assume normally distributed random effects) when the true distribution is different than the assumed.

However, this does not directly address the question of misspecification, which requires fixing the true distribution and varying the assumed distribution. To see why this is the case, consider a hypothetical situation where the standard deviation of an estimated parameter is 1 when the true and assumed random effects distributions are normal, but is equal to 2 when the true and assumed random effects distributions are t with 3 degrees of freedom. Further, suppose that, when the true distribution is t_3 , but we assume the distribution is normal, the standard deviation is 2.1. In this case, we would conclude that misspecification (assuming the distribution is normal when, in fact, it is t_3) has little impact, decreasing efficiency by 5%. But, if we compare the standard deviations under the two true distributions (1 under normality and 2.1 under t_3), we might be tempted to incorrectly conclude that misspecification causes a large decrease in efficiency. In Neuhaus, McCulloch and Boylan (2011), we therefore considered the same scenario as investigated in Litière, Alonso and Molenberghs (2007), but we varied the assumed distribution. Contrary to their conclusions, our results

showed that misspecification had virtually no impact on power, except a very modest impact for the combination of small sample sizes with huge random effects variances (variances of 16 or 32 on the logit scale). Litière, Alonso and Molenberghs (2008) consider the situation of random intercepts and slopes, which we consider in Section 12.

Similar results are true of investigations of informative cluster sizes with random intercepts. Hoffman, Sen and Weinberg (2001), Williamson, Datta and Satten (2003) and Benhin, Rao and Scott (2005) investigate the impact of informative cluster sizes on between-cluster covariates in logistic models via simulation. Hoffman, Sen and Weinberg (2001) states that, “We will show by a simulated scenario that when cluster size is nonignorable the behaviour of the generalised estimating equations approach breaks down.” and Williamson, Datta and Satten (2003) claim that, “... the usual generalized estimating equation approach resulted in severely biased estimates of both the marginal regression and association parameters.” These suggest a general failure of generalized estimating equation approaches. However, careful examination of their results shows that the poor performance is mainly isolated to the intercept terms. The simulation reported in Table 3 of Hoffman, Sen and Weinberg (2001) does not demonstrate bias for b_1 , the covariate effect, and simulations reported in Williamson, Datta and Satten (2003) show bias on the order of 5% (their Table 1) for covariate effects. Similar results hold for maximum likelihood and random intercept models (Neuhaus and McCulloch, 2011).

In summary, bias in estimates of between-cluster covariates has not been demonstrated. Efficiency loss has been demonstrated, but only in distributions quite far from non-normality, indicating a high degree of robustness. The degree of robustness is similar to that of normality-based tests of means.

6. ESTIMATION OF AN INTERCEPT

In many cases, linear mixed models assuming normality give unbiased or nearly unbiased estimates of all the fixed effect parameters including the intercept, even when the random intercepts distribution is non-normal. However, the same is not true for nonlinear models. Both theoretical and simulation studies have shown that estimates of the intercept may be biased when the random effects distribution is far from normal. Neuhaus, Hauck and Kalbfleisch (1992) give a local approximation for logistic models relating bias of

estimators of the intercept under assumed normality for the random intercepts to asymmetry of the random effects distribution. Heagerty and Kurland (2001) show asymptotic relative bias on the order of 20% or more when the true distribution is gamma but the assumed normal.

Similar results hold for informative cluster sizes. Williamson, Datta and Satten (2003) demonstrate bias of approximately 30% in estimating the intercept using independence generalized estimating equations using simulations for a logistic model. Similarly, Neuhaus and McCulloch (2011) demonstrate modest bias in estimating the intercept using maximum likelihood methods. They give a theoretical argument as to why bias is to be expected in the intercept but not other regression parameters.

In summary, misspecification of the shape of the random effects distribution can introduce a moderate to large bias in estimation of the intercept. This will carry over to possible bias in estimation of the mean value of the outcome for fixed covariate values. So, when inference focuses on mean estimation or the intercept, care should be taken with the specification of the random effects distribution.

7. PREDICTION OF THE RANDOM EFFECT, b_i

Much less work exists on the effect of misspecification on prediction of the random effects. Magder and Zeger (1996) demonstrated robustness of predictions as gauged by mean square error of prediction in linear mixed models and conclude, "Differences in the performances of the various methods in estimation of the random effects, b , were generally not very large." In simulating the performance of the posterior mean estimates in logistic models, Agresti, Caffo and Ohman-Strickland (2004) did not find a change in performance when the true distribution was varied between uniform, exponential and normal and the assumed distribution was normal. However, as noted above, this does not directly address the question of misspecification. They did find a number of situations in which assuming a normal distribution suffered a moderate loss of performance when the true distribution was a discrete two-point distribution. Zhang et al. (2008) (their Table 2) showed only modest differences in the mean square error of prediction of random slopes and intercepts when assuming normality and under a true log gamma distribution for a linear mixed model. McCulloch and Neuhaus (2011) investigated both linear mixed models and logistic mixed models under a variety of assumed and true distributions, both by theory

and simulation, and found only modest impacts to misspecification on the mean square error of prediction, using posterior mean predictions for the linear mixed models and posterior mode predictions for the logistic mixed models. Exceptions were when a distribution with limited support was assumed, but the true distribution had a wider range of support and for large random effects variances and large cluster sizes.

The *shape* of the distribution of the best predicted values is a different matter. A number of authors (Lesaffre and Molenberghs, 2001; Magder and Zeger, 1996; Zhang and Davidian, 2001) have demonstrated convincingly it may not reflect the true underlying shape of the distribution of the b_i , but instead the assumed distribution. We also demonstrate this in our example. So, although the *performance* of the best predicted values (as gauged by the mean square error of prediction) is robust, shapes of distributions of best predicted values are not.

To recap, the shape of the distribution of the best predicted values is highly sensitive to the assumed form of the distribution. This is unfortunate because it means that diagnostics based on the empirical distribution of the best predicted values (e.g., histograms or Q-Q plots) are unreliable. However, the performance of either posterior mean or posterior mode predicted values (as judged by overall mean square error of prediction) are generally robust across a wide variety of assumed distributions. Some loss of efficiency was evidenced with large cluster sizes and large random effects variances, in which definition of the correct random effects distribution is clearer.

8. ESTIMATE THE RANDOM EFFECTS VARIANCE

Again, with regard to estimation of the random effects variance, there is less research. Neuhaus, Hauck and Kalbfleisch (1992) and Heagerty and Kurland (2001) demonstrated little effect on estimation of the random effects variance, with asymptotic relative bias on the order of 15% or less. Agresti, Caffo and Ohman-Strickland (2004) demonstrated some situations, mostly with larger random effects variances and large cluster sizes, when there was efficiency loss with assuming a normal distribution, when the true distribution was a discrete, two-point distribution. Neuhaus and McCulloch (2011) show via simulation in the informative cluster size case that bias is slight.

As noted, this aspect of misspecification has not received the scrutiny that the other aspects above have, perhaps because inference centering on the random effects variance is somewhat less common. As a result,

the literature is far from definitive. Nevertheless, the basic picture is similar to that of covariate effects: estimates of the random effects variance appears relatively robust to misspecification of the random intercept distributional shape.

9. MISSPECIFICATION OF OTHER ASPECTS OF THE MODEL

We have mainly considered the situation in which only a single aspect of the model, the shape of the random intercepts distribution, has been misspecified. In practice, of course, all assumptions are violated to at least a minor degree. What is the impact of simultaneous misspecification of more than one aspect? There is little in the literature to report, but previous work does inform on two aspects. First, Neuhaus and McCulloch (2006) show that use of conditional likelihood methods and methods that separate covariates into between and within components eliminate or reduce bias for the within-cluster covariate in the situation where the random effects are associated with the covariates. The argument works by showing that estimation of the within-cluster covariate is essentially divorced from the random effects distribution when using these methods, also making it impervious to shape misspecification. Second, the arguments of Neuhaus and McCulloch (2011), showing that informative cluster sizes have little impact on the bias of regression coefficients, proceed by converting the informative cluster size problem into the misspecified random effects distribution as displayed in (2). This same argument can be used to show that, even with simultaneous misspecification of the shape of the random effects distribution and informative cluster sizes, maximum likelihood estimators of the regression coefficients (other than the intercept) will be consistent in a linear mixed model and, for $\beta = 0$, consistent at zero for generalized linear mixed models. While not all-encompassing, this result for generalized linear mixed models is important for testing the null hypothesis that $\beta = 0$.

10. A SIMULATION STUDY

We performed a simulation study to evaluate the performance of each of the above aspects of inference. The goal was to gauge the performance of assumed normal fits to a distribution that was highly non-normal, but not as extreme as a two-point, discrete distribution. We chose a Tukey(g, h) distribution, since, depending on the values of g and h , the Tukey distribution can be

quite skewed and/or heavy-tailed. See He and Raghunathan (2006) for a recent reference. We chose $g = 0.5$ and $h = 0.1$, which gives a mean of 0.31, variance of 2.27, skewness of 3.41 and a kurtosis of 44.24.

The simulations used two covariates: one within-cluster and one between-cluster covariate. The within-cluster covariate was equally spaced between 0 and 1. The between-cluster covariate was binary with a 25%/75% division. The parameter values were: $\beta_0 = -2.5$, $\beta_{\text{between}} = 2$, $\beta_{\text{within}} = 1$, $\sigma_b = 1$. We set the number of clusters, m , to 200 and used a variety of cluster sizes ($n = 2, 4, 6, 10, 20$ and 40). We simulated data under a logistic link and the Tukey(g, h) distribution. We ran 1,000 replications for each scenario and used common random numbers across different cluster sizes and fitted distributions to increase precision of comparisons. We conducted simulations in SAS (Ver 9.1, SAS Institute, Cary NC) and fit models using Proc NLMIXED.

To each simulated data set we fit three GLMMs with a logistic link. The first model assumed that the random effects were standard normal and the second assumed they followed a standardized Tukey(g, h) distribution with unknown g and h . Especially with small cluster sizes, there is little or no information about the shape parameters of the Tukey distribution and trying to estimate those parameters led to unstable estimation. To directly assess the effect of misspecification and to avoid confounding the effects of an incorrect distributional shape with the estimation of the two additional parameters, we fit a third GLMM with the random effects following a Tukey distribution with g and h fixed at their true values of 0.5 and 0.1, but still estimating the random effects variance.

Figure 1 gives the results for the bias of the estimators. As expected, there was virtually no impact of using an assumed normal distribution in the estimation of the within-cluster covariate. There was a modest impact in estimating the between covariate and the intercept (but with bias less than 5% for the between-cluster covariate and less than 10% for the intercept). Impact on estimation of the log standard deviation of the random effects was negligible.

As noted, estimation of the g and h parameters for the Tukey distribution introduced instability in estimating the log standard deviation of the random effects. This led to poor estimation of the intercept for cluster size $n = 2$ and for the log standard deviation for cluster sizes $n = 2, 4, 6$ and 10. Perhaps surprisingly, the estimation of g and h had little impact on estimation of β_{between} .

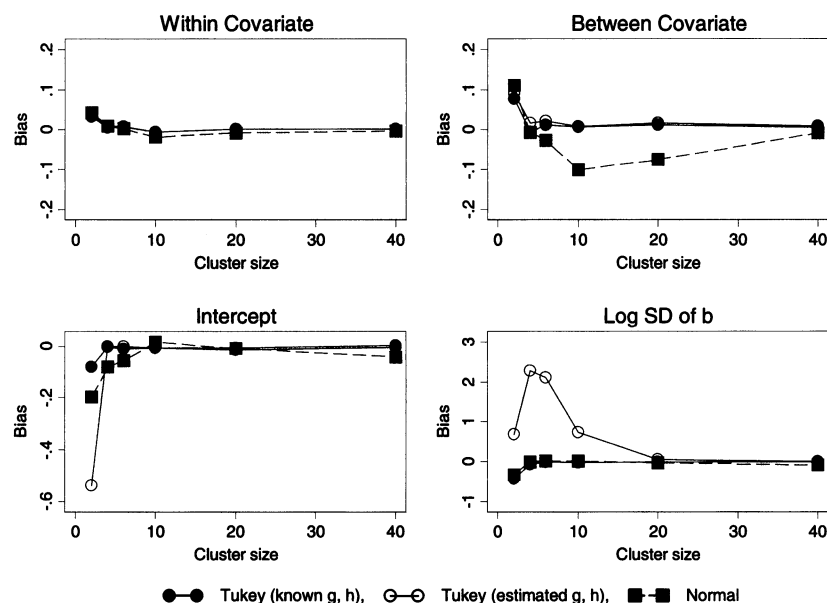


FIG. 1. Bias of estimators of the parameters of a logistic mixed effects model, (1), with Tukey(g, h) distributed random intercepts, fit with assumed normal and Tukey distributions.

Fitting these models is quite involved, so we give some further details on the simulation. Not all the replications in the simulations achieved convergence via NLMIXED. Fitting the assumed normal and Tukey distribution with known g and h had excellent convergence rates, all above 95%. However, the Tukey distribution with estimated g and h was numerically challenging and convergence rates for $n = 4$ and 6 were lower (82.5% and 83.8%, respectively), with estimation of $\log \sigma_b$ especially problematic. The figures give results for all the runs, using final parameter values in cases with nonconvergence. Comparing the convergent runs only with all the runs showed little impact on the bias of the parameters in Figure 1, excepting $\log \sigma_b$. Simulation code is available from the authors.

Figure 2 gives the results for the standard deviations of the estimators. The standard deviations of the estimators were nearly the same under both the fitted normal and Tukey distributions across all the cluster sizes and for all the parameters, with the exception being a modest loss of efficiency for the between-cluster covariate, especially with larger cluster sizes. So misspecifying the random effects distribution produced essentially no loss in estimation efficiency.

Since estimation of between-cluster covariates is an important inferential goal in clustered data settings and there is a modest impact of misspecifying the distribution as normal, we give more detail on the actual distributions of the estimates of the between-cluster covariate effect. Figure 3 shows boxplots of the estimates

under the assumed normal and Tukey (with fixed g and h) fits. Behavior of the assumed normal fit is slightly worse than the true Tukey fit for cluster sizes of 10 and 20 and somewhat more variable for a cluster size of 40. For smaller cluster sizes the results are comparable. Given the extreme differences between a normal and Tukey(0.5, 0.1) distribution, this represents a large degree of robustness, especially for cluster sizes 6 or smaller.

SAS Proc NLMIXED calculated best predicted values as modes of the log of $f(\mathbf{y}_i | \mathbf{x}_i, b_i) f_b(b_i | \sigma_b)$. Figure 4 shows the mean squared error of prediction under the three fitted models. Compared to the Tukey distribution with estimated g and h , the misspecified normal actually outperforms it up until cluster sizes of 10. Compared to the Tukey distribution with fixed g and h , the performance is slightly worse, with the inflation in mean square error of prediction ranging from about 5% at $n = 2$ to about 20% for the larger cluster sizes. Again, given the extreme differences between a normal and Tukey, this is a high degree of robustness.

These results support the broad conclusion that, for distributional shapes quite different from the normal, estimation of the intercept may be biased. However, the bias for estimating other parameters was low and efficiency high. For prediction of random effects, the mean square error of prediction was modestly increased when incorrectly assuming normality, especially for larger cluster sizes.

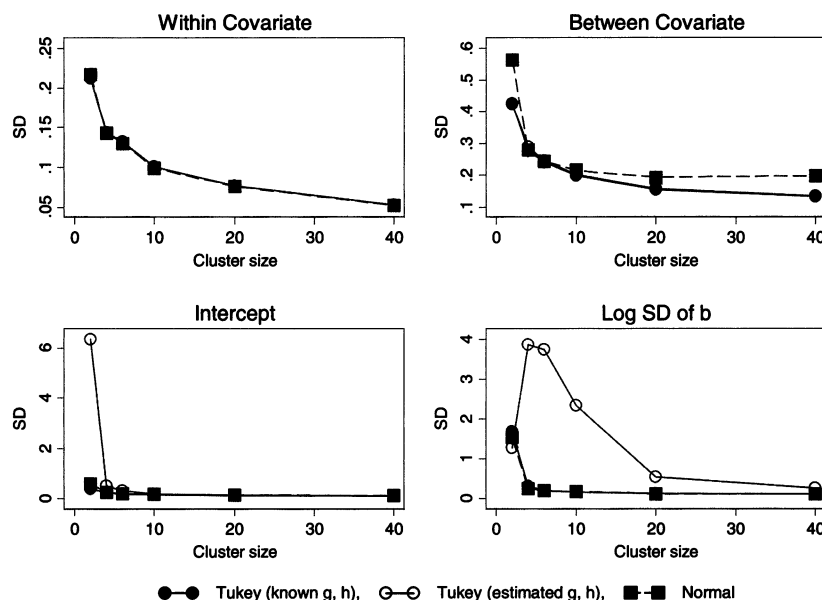


FIG. 2. Standard deviation of estimators of the parameters of a logistic mixed effects model, (1), with Tukey(g, h) distributed random intercepts, fit with assumed normal and Tukey distributions.

11. EXAMPLE—HERS

HERS was a randomized, blinded, placebo controlled trial for women with previous coronary disease. The study enrolled 2,763 women and followed them annually for five subsequent visits. We will consider only the subset of $N = 1,378$ that were not diabetic and who had systolic blood pressure less than 140 at the beginning of the study. We treat HERS as a prospective cohort study using the first four visits. We modeled the binary outcome of high blood pressure (HBP), defined

as a systolic blood pressure greater than 140, as a function of visit (numerical 0 through 3), body mass index (BMI) and whether or not the participant was on high blood pressure medication at that visit (HTN). The visit trend variable is essentially a within-person covariate and BMI and HTN are mostly between-person covariates (with 95% and 76%, respectively, of their variability associated with between-person differences).

With HBP_{it} equal to 1 if woman i at visit t had high blood pressure (and 0 otherwise), our random intercepts logistic regression model was

$$\begin{aligned} \text{logit}(P\{HBP_{it} = 1\}) \\ = \beta_0 + b_{0i} + \beta_1 \text{visit}_t + \beta_2 \text{BMI}_{it} + \beta_3 \text{HTN}_{it}, \end{aligned} \quad (4)$$

where $b_{0i} \sim \text{i.i.d. } \mathcal{N}(0, \sigma_b^2)$ or
 $b_{0i} \sim \text{i.i.d. } \sigma_b \{\text{Exp}(1) - 1\}$ or
 $b_{0i} \sim \text{i.i.d. discrete with three mass points}$ or
 $b_{0i} \sim \text{i.i.d. } \sigma_b \{\text{Tukey}(g, h)\}.$




FIG. 3. Boxplots of estimates of β_{between} from a logistic mixed effects model, (1), with Tukey(g, h) distributed random intercepts, fit with assumed normal and Tukey distributions. Reference line given at the true value of $\beta_{\text{between}} = 2$.

We fit models via maximum likelihood to the data using, in turn, each of the assumed random effects distributions given in (4). We used SAS Proc NLMIXED (SAS Institute, Cary, NC) for the continuous distributions, which gives posterior mode estimates of the b_{0i} . We used the Stata module GLLAMM (www.gllamm.org) to fit the three point discrete distribution, which gives posterior mean estimates of the b_{0i} .




FIG. 4. Mean square error of prediction of the random effects from a logistic mixed effects model, (1), with Tukey(g, h) distributed random intercepts, fit with assumed normal and Tukey distributions.

We chose the exponential (Exp) and Tukey distributions since they are parametric but quite different from the normal. The exponential distribution has scale parameter 1, a bounded range, and is skewed right. The Tukey distribution is a flexible parametric family that can take on a wide variety of shapes (He and Raghunathan, 2006); it includes the normal as a special case ($g = h = 0$) but also includes distributions with extreme skewness and kurtosis. The discrete distribution is similar to a nonparametric maximum likelihood fit, but specifies a priori the number of mass points.

Table 1 lists the fitted coefficients and maximized log likelihood values. As expected, the fixed effects parameter estimates are quite similar, especially the “within” covariate, Visit, even though there are modest differences in the fits of the models as judged by the value of the maximized log likelihood. The shape parameters for the Tukey distribution had large standard errors but were estimated to be $g = 2.5$ and $h = -1.8$, which is a distribution with bounded range, and is slightly

skewed right and heavy tailed (skewness of about 0.9 and kurtosis of about 2.3).

Figure 5 gives histograms of the best predicted values of the random intercept deviations, b_{0i} , under each of the above assumed distributions. As is evident from the figure, the shape of the distribution of the best predicted values is dependent on the assumed random effects distribution, with the discrete and exponential distributions especially having different shapes than the normal and Tukey, which are similar to one another.

12. RANDOM INTERCEPTS AND SLOPES

The remarks above pertain mostly to models with random intercepts only. An understudied area is what happens with the further complication of random coefficient models, such as random intercepts and slopes. The line of argument cited above using minimization of Kullback–Leibler divergence (Neuhaus, Hauck and Kalbfleisch, 1992; Neuhaus, Kalbfleisch and Hauck, 1994; Heagerty and Kurland, 2001) can also be used in the random intercepts and slopes situation. This line of argument shows, for the linear mixed model, that estimates of fixed effects parameters will be consistent. However, for the linear mixed model this is well known (Verbeke and Lesaffre, 1997) by alternate means. For generalized linear mixed models, results analogous to the random intercepts case hold. For fixed effects for which there is a corresponding, misspecified random effects distribution, bias can result. But fixed effects orthogonal to random effects are little affected. For example, if a covariate is both a fixed effect and a random effect (i.e., a random slope), the estimate of the fixed effect for that covariate may exhibit bias when the distribution of the random slope is misspecified. But estimates of covariate effects, orthogonal to the one which is random, are little affected. For the informative cluster size situation, Neuhaus and McCulloch (2011) make these arguments more precise. In the case

TABLE 1
HERS model fit comparisons with different assumed random effect distributions. Model-based standard errors for the fixed effects are given as subscripts

Random effects distribution	–2 loglik	Parameter estimates				
		Intercept	Visit	BMI	HTN	$\log(\hat{\sigma}_b)$
Normal	3695	–4.28 _{0.41}	0.86 _{0.05}	0.024 _{0.013}	–0.36 _{0.16}	0.49
Exponential	3732	–3.96 _{0.38}	0.84 _{0.05}	0.023 _{0.013}	–0.30 _{0.14}	0.19
Discrete	3674	–4.05 _{0.38}	0.87 _{0.05}	0.021 _{0.012}	–0.37 _{0.16}	0.23
Tukey	3677	–4.10 _{1.06}	0.86 _{0.05}	0.022 _{0.013}	–0.36 _{0.16}	0.28

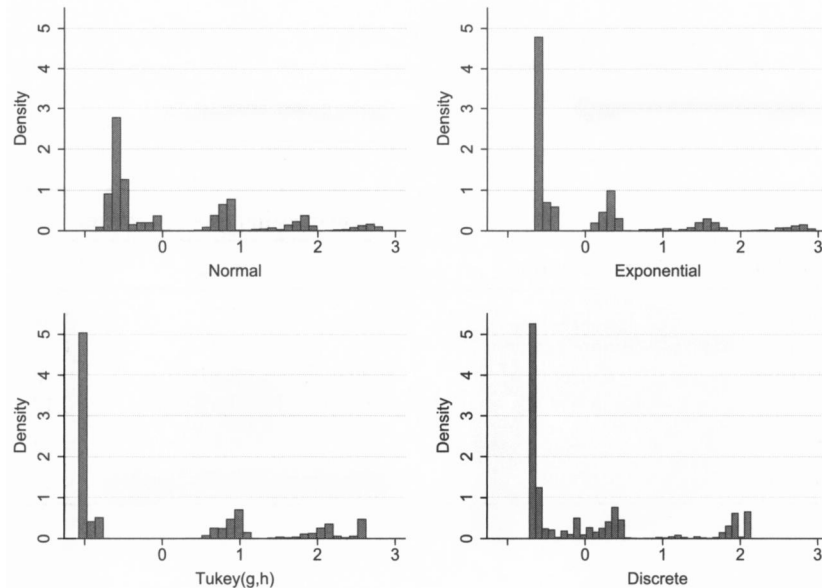


FIG. 5. Best predicted values for the HERS data under different assumed random effects distributions.

of prediction of realized values or random effects, McCulloch and Neuhaus (2011) report limited results for random intercepts and slopes; they show a fair degree of robustness for mean square error of prediction and comparable behavior to the random intercepts case.

Contrary to the above, Litière, Alonso and Molenberghs (2008) describe some simulations for the random intercepts and slopes situation that report to show significant bias for a mixed effects logistic model with true normal or true mixture of normal distributions to which an assumed bivariate normal distribution is fit. The situation they simulate is quite extreme. The true bivariate mixture distribution corresponding to their least extreme configuration, which they call V_1 , is a mixture of standard normals with mean values at plus or minus 2 for both the random intercept and slope; this generates a distribution (in two dimensions) with two isolated peaks. The standard deviation on the logit scale for both the intercepts and slopes for this model is $\sqrt{5} \approx 2.24$.

For the random intercept, a random effect one standard deviation above its mean of 0 would have an odds of the outcome that is $\exp(2.24) \approx 9.36$ higher than a random effect at its mean of 0. So this is a moderately large effect. The effect of a random slope one standard deviation above its mean of 0 depends on the values of the associated covariate, which range from 0 to 8 in their simulation. So a random slope one standard deviation above its mean of 0 would have the same odds of the outcome at $t = 0$ compared to an “average” random slope, but at $t = 8$ would have an odds of

the outcome of $\exp(8\sqrt{5}) \approx 58,700,000$ times higher. Furthermore, the correlation between the random intercepts and slopes was 0.9. The other scenarios they simulate are even more extreme! A more reasonable value for the random slope variance would be approximately 0.08. At $t = 8$ this would generate an odds of an outcome which is $\exp(8\sqrt{0.08}) \approx 9.61$ times higher, an effect comparable in magnitude to the random intercepts.

They also report poor convergence rates when fitting the assumed bivariate normal distribution to the bivariate mixture distribution. We performed a small simulation to check their results and to assess the magnitude of the bias under the highly non-normal but more reasonable situation where the random slopes variance was 0.08. We used the same model as Litière, Alonso and Molenberghs (2008), namely,

$$(5) \quad \begin{aligned} & \text{logit}(P\{Y_{it} = 1\}) \\ &= (\beta_0 + b_{0i}) + \beta_b z_i + (\beta_w + b_{wi})t_j, \end{aligned}$$

with z_i binary with equal proportions of 0 and 1, $t_j = 0, 1, 2, 4, 6$, and 8 (within each cluster of size 6) and 100 clusters. We performed 500 replications and fit the models via maximum likelihood using SAS Proc NLMIXED assuming a bivariate normal distribution. Table 2 reports the results. When fitting to a true normal distribution, the median values when the random slope variance is 5 (upper panel) are remarkably close to the true values. When fitting to a true mixture distribution, the median values exhibit bias, but not nearly

TABLE 2

Medians of maximum likelihood estimates of fixed effects from fitting model (5) assuming a bivariate normal distribution when the true distribution is normal or a mixture of bivariate normals. β_0 is the intercept, β_b is the between-cluster coefficient and β_w is the within-cluster coefficient. In each case the random intercepts had variance 5 and a correlation with the random slopes of 0.9. The random slopes variance was either 5 (upper panel) or the more reasonable value of 0.08 (lower panel). The simulation generated 500 replications with 100 clusters of size 6

	True dist'n	Simulation	Median estimates (true values)			Convergence rate
			β_0 (-6)	β_b (2)	β_w (1)	
$\text{var}(b_{wi}) = 5$	Normal	Current	-6.26	2.13	1.07	98%
	Normal	Litiere et al.	-6.14	2.04	1.04	not given
	Mixture	Current	-6.34	2.87	0.86	99%
	Mixture	Litiere et al.	-10.52	2.52	-0.13	<43%
$\text{var}(b_{wi}) = 0.08$	Normal	Current	-6.18	2.12	1.04	100%
	Mixture	Current	-5.89	2.26	0.95	100%

as extreme as those reported by Litière, Alonso and Molenberghs (2008). Convergence rates when fitting to the mixture distribution are much better as well. Our results are consistent with previous literature, in that the impact on the within-cluster coefficient β_w is less extreme than the between-cluster coefficient, β_b . Under the more reasonable random slopes variance of 0.08 there is only a small amount of bias. Of course, it is important to remember that, even under the more reasonable value of the variance, the mixture model is a highly non-normal bivariate distribution with most of its mass concentrated around $(-2, -2)$ and $(2, 2)$.

13. SUMMARY

We considered robustness to the assumed distribution for a random intercept when using maximum likelihood methods for fitting a generalized linear mixed model such as (1). In practice, data analysts commonly assume normality for the random effects. Theory and simulation studies indicate that most aspects of statistical inference are highly robust to this assumption. Especially robust were inferences for within-cluster covariate effects, which are often a key reason for considering a clustered data design.

These conclusions are contrary to much of the previous literature. We have argued that this is because (a) results for the nonclustered data situation (e.g., Heckman and Singer, 1984) are incorrectly interpreted as relevant to the clustered data setting, (b) results restricted to one portion of the model (primarily the intercept) are characterized as a general failure of the methodology, and (c) reanalysis of previous scenarios has sometimes given different interpretations or results.

Exceptions to this robustness include estimation of the intercept, which can be biased when the random intercept distribution is misspecified. The shape of the estimated random effects distribution (based on the distributional shape of the predicted random effects) can also be quite sensitive to the shape of the assumed distribution and thus not reflect the true distribution.

However, a wide array of inferences are quite robust to this type of misspecification, including estimation of covariate effects, estimation of the random effects variance and the mean square error of prediction of the realized value of random effects. We also argued that this robustness extends to random intercept situations where the cluster size is informative.

Between-cluster covariate effects and estimation of the log of the standard deviation of the random effects were found to be robust, but not impervious, to misspecification of the random effects distribution. The literature and results here indicate that a modest amount of bias and minor loss of efficiency can occur, especially when the true distribution is far from the assumed (e.g., assuming a normal distribution when the true distribution is a two-point distribution or the Tukey distribution investigated here), the random effects are large, and the cluster size is large. When random effects and cluster sizes are large, the data provide a fair amount of information as to the shape of the random effects distribution and correctly specifying the random effects distribution gives some advantage.

Particularly robust was estimation of within-cluster covariate effects. This is especially important because exploitation of within-cluster comparisons is often the rationale for conducting longitudinal or clustered data studies. Virtually all the evidence in the literature and

the new results reported here indicate that within-cluster covariate effects are estimated with no more bias when fitting a misspecified model than when fitting the correct model and with little or no loss of efficiency.

ACKNOWLEDGMENTS

We thank Ross Boylan for computational assistance with the simulation studies and Stephen Hulley for use of the HERS data set. Support was provided by NIH Grant R01 CA82370.

REFERENCES

- AGRESTI, A., CAFFO, B. and OHMAN-STRICKLAND, P. (2004). Examples in which misspecification of a random effects distribution reduces efficiency, and possible remedies. *J. Comput. Graph. Statist.* **47** 639–653. MR2100566
- AITKIN, M. (1999). A general maximum likelihood analysis of variance components in generalized linear models. *Biometrics* **55** 117–128. MR1705676
- AKAIKE, H. (1973). Information theory and an extension of the maximum likelihood principle. In *Second International Symposium on Information Theory* (B. N. Petrov and F. Czaki, eds.) 267–281. Akademiai Kiadó, Budapest. MR0483125
- AUBLE, T. E., HSIEH, M., MCCAUSLAND, J. B. and YEALY, D. M. (2007). Comparison of four clinical prediction rules for estimating risk in heart failure. *Annals of Emergency Medicine* **50** 127–135.
- BENHIN, E., RAO, J. N. K. and SCOTT, A. J. (2005). Mean estimating equation approach to analysing cluster-correlated data with nonignorable cluster sizes. *Biometrika* **92** 435–450. MR2201369
- BUTLER, S. M. and LOUIS, T. A. (1992). Random effects models with non-parametric priors. *Stat. Med.* **11** 1981–2000.
- CAFFO, B., AN, M. and ROHDE, C. (2007). Flexible random intercept models for binary outcomes using mixtures of normals. *Comput. Statist. Data Anal.* **51** 5220–5235. MR2370867
- CHEN, J., ZHANG, D. and DAVIDIAN, M. (2002). A Monte Carlo EM algorithm for generalized linear mixed models with flexible random effects distribution. *Biostatistics (Oxford)* **3** 347–360.
- DAVIDIAN, M. and GALLANT, A. R. (1993). The nonlinear mixed effects model with a smooth random effects density. *Biometrika* **80** 475–488. MR1248015
- GHIHEY, W., LESAFFRE, E. and EILERS, P. (2004). Smooth random effects distribution in a linear mixed model. *Biometrics* **60** 945–953. MR2133547
- GHIHEY, W., LESAFFRE, E. and VERBEKE, G. (2010). A comparison of methods for estimating the random effects distribution of a linear mixed model. *Stat. Methods Med. Res.* **19** 565–600.
- HE, Y. and RAGHUNATHAN, T. (2006). Tukey's gh distribution for multiple imputation. *Amer. Statist.* **60** 251–256. MR2246758
- HEAGERTY, P. J. and KURLAND, B. F. (2001). Misspecified maximum likelihood estimates and generalised linear mixed models. *Biometrika* **88** 973–985. MR1872214
- HEAGERTY, P. J. and ZEGER, S. L. (2000). Marginalized multi-level models and likelihood inference (with comments and a rejoinder by the authors). *Statist. Sci.* **15** 1–26. MR1842235
- HECKMAN, J. and SINGER, B. (1984). A method for minimizing the impact of distributional assumptions in econometric models for duration data. *Econometrica* **52** 271–320. MR0735309
- HOFFMAN, E. B., SEN, P. K. and WEINBERG, C. R. (2001). Within-cluster resampling. *Biometrika* **88** 1121–1134. MR1872223
- HUANG, X. (2009). Diagnosis of random-effect model misspecification in generalized linear mixed models for binary response. *Biometrics* **65** 361–368. MR2751459
- HUBER, P. J. (1967). The Behavior of Maximum Likelihood Estimates Under Nonstandard Conditions. In *Proc. Fifth Berkeley Sympos. Math. Statist. and Probability* (Berkeley, Calif., 1965/66) **1** (L. M. Le Cam and J. Neyman, eds.) 221–223. Univ. California Press, Berkeley. MR0216620
- HULLEY, S., GRADY, D., BUSH, T., FURBERG, C., HERRINGTON, D., RIGGS, B. and VITTINGHOFF, E. (1998). Randomized trial of estrogen plus progestin for secondary prevention of coronary heart disease in postmenopausal women. Heart and Estrogen/progestin Replacement Study (HERS) Research Group. *Journal of the American Medical Association* **280** 605–613.
- KULLBACK, S. C. (1959). *Information Theory and Statistics*. Wiley, New York. MR0103557
- LAIRD, N. (1978). Nonparametric maximum likelihood estimation of a mixing distribution. *J. Amer. Statist. Assoc.* **73** 805–811. MR0521328
- LEE, Y. and NELDER, J. A. (2004). Conditional and marginal models: Another view. *Statist. Sci.* **19** 219–238. MR2140539
- LESAFFRE, E. and MOLENBERGHS, G. (2001). Multivariate probit analysis: A neglected procedure in medical statistics. *Stat. Med.* **10** 1391–1403.
- LITIÈRE, S., ALONSO, A. and MOLENBERGHS, G. (2007). Type I and type II error under random-effects misspecification in generalized linear mixed models. *Biometrics* **63** 1038–1044. MR2414580
- LITIÈRE, S., ALONSO, A. and MOLENBERGHS, G. (2008). The impact of a misspecified random-effects distribution on the estimation and the performance of inferential procedures in generalized linear mixed models. *Stat. Med.* **27** 3125–3144. MR2522153
- MAGDER, L. S. and ZEGER, S. L. (1996). A smooth nonparametric estimate of a mixing distribution using mixtures of Gaussians. *J. Amer. Statist. Assoc.* **91** 1141–1151. MR1424614
- MCCULLOCH, C. E. and NEUHAUS, J. M. (2011). Prediction of random effects in linear and generalized linear models under model misspecification. *Biometrics* **67** 270–279.
- MCCULLOCH, C. E., SEARLE, S. R. and NEUHAUS, J. M. (2008). *Generalized, Linear and Mixed Models*, 2nd ed. Wiley, New York. MR2431553
- METLAY, J. P., CAMARGO, C. A., MACKENZIE, T., MCCULLOCH, C. E., MASELLI, J., LEVIN, S. K., KERSEY, A., GONZALES, R. and THE IMPAACT INVESTIGATORS (2007). Cluster-randomized trial to improve antibiotic use for adults with acute respiratory infections treated in emergency departments. *Annals of Emergency Medicine* **50** 221–230.
- NEUHAUS, J. M., HAUCK, W. W. and KALBFLEISCH, J. D. (1992). The effects of mixture distribution misspecification when fitting mixed-effects logistic models. *Biometrika* **79** 755–762.

- NEUHAUS, J. M., KALBFLEISCH, J. D. and HAUCK, W. W. (1994). Conditions for consistent estimation in mixed-effects models for binary matched pairs data. *Canad. J. Statist.* **22** 139–148. MR1271451
- NEUHAUS, J. M. and KALBFLEISCH, J. D. (1998). Between- and within-cluster covariate effects in the analysis of clustered data. *Biometrics* **54** 638–645.
- NEUHAUS, J. M. and MCCULLOCH, C. E. (2006). Separating between and within-cluster covariate effects using conditional and partitioning methods. *J. Roy. Statist. Soc. Ser. B* **68** 859–872. MR2301298
- NEUHAUS, J. M., MCCULLOCH, C. E. and BOYLAN, R. (2011). A note on type II error under random effects misspecification in generalized linear mixed models. *Biometrics* **67** 654–656.
- NEUHAUS, J. M. and MCCULLOCH, C. E. (2011). Estimation of covariate effects in generalised linear mixed models with informative cluster sizes. *Biometrika* **98** 147–162.
- PIEPHO, H.-P. and MCCULLOCH, C. E. (2004). Transformations in mixed models: Application to risk analysis for a multi-environment trial. *J. Agric. Biol. Environ. Statist.* **9** 123–137.
- RASCH, D. and GUIARD, V. (2004). The robustness of parametric statistical methods. *Psychology Science* **46** 175–208.
- RAUDENBUSH, S. and BRYK, A. (2002). *Hierarchical Linear Models: Applications and Data Analysis Methods*, 2nd ed. Sage Publications, Thousand Oaks.
- SEARLE, S. R., CASELLA, G. and MCCULLOCH, C. E. (1992). *Variance Components*. Wiley, New York. MR1190470
- SELBY, J. V., FIREMAN, B. H., LUNDSTROM, R. J., SWAIN, B. E., TRUMAN, A. F., WONG, C. C., FROELICHER, E. S., BARRON, H. V. and HLATKY, M. A. (1996). Variation among hospitals in coronary-angiography practices and outcomes after myocardial infarction in a large health maintenance organization. *New England Journal of Medicine* **335** 1888–1896.
- TAO, H., PALTA, M., YANDELL, B. S. and NEWTON, M. A. (1999). An estimation method for the semiparametric mixed effects model. *Biometrics* **55** 102–110. MR1705675
- VERBECKE, G. and LESAFFRE, E. (1997). The effect of misspecifying the random-effects distribution in linear mixed models for longitudinal data. *Comput. Statist. Data Anal.* **23** 541–556. MR1437679
- WHITE, H. (1994). *Estimation, Inference, and Specification Analysis*. Cambridge Univ. Press, Cambridge. MR1292251
- WILLIAMSON, J. M., DATTA, S. and SATTEN, G. A. (2003). Marginal analyses of clustered data when cluster size is informative. *Biometrics* **59** 36–42. MR1978471
- ZHANG, D. and DAVIDIAN, M. (2001). Linear mixed models with flexible distribution of random effects for longitudinal data. *Biometrics* **57** 795–802. MR1859815
- ZHANG, P., SONG, P. X. K., QU, A. and GREENE, T. (2008). Efficient estimation for patient-specific rates of disease progression using nonnormal linear mixed models. *Biometrics* **64** 29–38. MR2422816